

# Understanding Indexing in MSSQL

## What is Indexing in MSSQL?

Indexing in MSSQL (Microsoft SQL Server) is a database optimization technique used to improve the speed of data retrieval. It creates a data structure (similar to the index in a book) that helps the database engine locate rows more efficiently.

### Types of Indexes in MSSQL

1. **Clustered Index:**

- A clustered index determines the physical order of data in the table.
- There can only be one clustered index per table since the data can only be sorted in one way.

2. **Example:**

```
CREATE CLUSTERED INDEX IX_Products_ProductID ON Products(ProductID);
```

### Non-Clustered Index:

- A non-clustered index creates a separate structure from the table data, with pointers to the actual rows.
- Multiple non-clustered indexes can be created on a single table.

**Example:**

```
CREATE NONCLUSTERED INDEX IX_Products_ProductName ON  
Products(ProductName);
```

## Full-Text Index:

- Used for advanced text-based searches, such as searching for keywords in large text columns.

## Example:

```
CREATE FULLTEXT INDEX ON Products(ProductName)
```

```
KEY INDEX IX_Products_ProductID;
```

## Benefits of Indexing

1. **Faster Query Performance:**
  - Indexes speed up queries by reducing the amount of data scanned.
2. **Efficient Sorting:**
  - Queries with **ORDER BY** clauses benefit from indexes.
3. **Improved Filtering:**
  - **WHERE** clauses can quickly locate rows matching criteria.

# What Happens When You Create an Index?

## Clustered Index:

- It **reorganizes the actual table data** to match the index order. The table itself is stored in the sorted order of the indexed column(s).
- Think of it as sorting the rows of the table by the indexed column (like sorting a book alphabetically by title).

## Non-Clustered Index:

- It creates a **separate structure** (like a lookup table) that stores the indexed column(s) and pointers to the actual rows in the table.
- The table data remains unchanged, but the index provides a faster way to locate rows based on the indexed column.

## 2. Does It Create a JSON File or Table?

- **No JSON File:**

- MSSQL does not use JSON files for indexing. The index is stored in the database in its own internal structure.

- **No New Table:**

- A new table is not created. The index is a special data structure linked to the table it belongs to.

Instead, the index is stored as **metadata** and **data pages** in the database files (**.mdf** and **.ndf** files).

### 3. What Does the Index Look Like Internally?

Indexes are typically implemented as a **B-tree** structure in MSSQL:

- **Clustered Index:**
  - The table data itself is stored in the B-tree structure.
  - Each leaf node contains the actual table data.
- **Non-Clustered Index:**
  - The B-tree contains pointers (references) to the actual data in the table.

#### Example of a Non-Clustered Index:

If you create an index on `ProductName`, the index might look like this internally:

Index:

- Apple → Pointer to Row 3
- Banana → Pointer to Row 1
- Cherry → Pointer to Row 2

The actual table remains unchanged:

Table:

Row 1: Banana

Row 2: Cherry

Row 3: Apple

## 4. Where Is the Index Stored?

- Indexes are stored in the same **database files** (.mdf or .ndf).
- When you query or rebuild an index, SQL Server uses these files to manage the index.

## 5. What Does It Mean for Storage?

- Indexes take up **additional storage space** within the database.
  - Clustered Index: Reorganizes the table, so no additional space is required for the index itself (but sorting may increase fragmentation).
  - Non-Clustered Index: Requires extra space for the index structure since it is stored separately from the table.

When you create an index in MSSQL:

- It creates an **internal data structure** (like a B-tree).
- No new table or JSON file is created.
- The index is stored within the database files and works alongside the table to improve query performance.



## Challenges of Indexing

1. **Performance Overhead:**
  - Indexes slow down **INSERT**, **UPDATE**, and **DELETE** operations because the index also needs to be updated.
2. **Storage Costs:**
  - Indexes consume additional disk space.
3. **Maintenance:**
  - Indexes can become fragmented and require periodic rebuilding.

## Part 2: Introduction to Elasticsearch

### What is Elasticsearch?

Elasticsearch is a distributed, open-source search and analytics engine designed for scalability and real-time data access. Unlike MSSQL, which is optimized for structured data, Elasticsearch excels in:

- **Full-text search:** Finding documents based on keywords and relevance.
- **Unstructured or semi-structured data:** JSON-based storage.
- **Distributed Architecture:** Handles petabytes of data across clusters.

# Why Use Elasticsearch?

## 1. **Full-Text Search:**

- Advanced capabilities like fuzzy search, autocomplete, and relevance scoring.

## 2. **Real-Time Analytics:**

- Perform aggregations and queries on massive datasets.

## 3. **Scalability:**

- Elastic scales horizontally by adding nodes to the cluster.

## 4. **JSON Document Storage:**

- Stores data in a schema-free JSON format.

# How Elasticsearch Works

## 1. **Index:**

- Similar to a database in SQL, an Elasticsearch index is a collection of documents.

## 2. **Document:**

- Each record (stored as JSON) in an Elasticsearch index is a document.

## 3. **Shard:**

- Indexes are divided into smaller units called shards for distribution.

## 4. **REST API:**

- Elasticsearch provides RESTful APIs for data ingestion, search, and analytics.

## Part 3: Using MSSQL and Elasticsearch in Docker

### Step 1: Running MSSQL in Docker

1. Pull and run the MSSQL Docker image:

```
docker pull mcr.microsoft.com/mssql/server:2019-latest
```

```
docker run -e 'ACCEPT_EULA=Y' -e 'SA_PASSWORD=YourStrongPassword123'  
-p 1433:1433 --name mssql-container -d  
mcr.microsoft.com/mssql/server:2019-latest
```

Connect to MSSQL and create the **Products** table:

```
CREATE DATABASE ProductDB;
```

```
USE ProductDB;
```

```
CREATE TABLE Products (
```

```
    ProductID INT PRIMARY KEY,
```

```
    ProductName NVARCHAR(100),
```

```
    Price DECIMAL(10, 2),
```

```
    Stock INT
```

```
);
```

```
INSERT INTO Products (ProductID, ProductName, Price, Stock) VALUES
```

```
(1, 'Laptop', 1200.00, 50),
```

```
(2, 'Mouse', 25.00, 200),
```

```
(3, 'Keyboard', 45.00, 100);
```

Add an index to the `ProductName` column:

```
CREATE NONCLUSTERED INDEX IX_Products_ProductName ON Products(ProductName);
```

## Step 2: Running Elasticsearch in Docker

1. Pull and run the Elasticsearch Docker image:

```
docker pull docker.elastic.co/elasticsearch/elasticsearch:8.10.0
```

```
docker run -d --name elasticsearch -p 9200:9200 -e "discovery.type=single-node" docker.elastic.co/elasticsearch/elasticsearch:8.10.0
```

Verify Elasticsearch is running:

- Open a browser and go to <http://localhost:9200>.

### Step 3: Integrating MSSQL and Elasticsearch in .NET Web API

#### 1. Create a .NET Web API Project

1. Create a new Web API project:

```
dotnet new webapi -n ProductSearchAPI
```

```
cd ProductSearchAPI
```

Install required NuGet packages:

```
dotnet add package Microsoft.EntityFrameworkCore
```

```
dotnet add package Microsoft.EntityFrameworkCore.SqlServer
```

```
dotnet add package NEST
```

## 2. Configure MSSQL

1. Add the MSSQL connection string to `appsettings.json`:

```
{  
  "ConnectionStrings": {  
    "DefaultConnection": "Server=localhost,1433;Database=ProductDB;User  
Id=SA;Password=YourStrongPassword123;"  
  },  
  "Elasticsearch": {  
    "Uri": "http://localhost:9200"  
  }  
}
```



Create the **Product** model and database context:

```
public class Product
```

```
{
```

```
    public int ProductID { get; set; }
```

```
    public string ProductName { get; set; }
```

```
    public decimal Price { get; set; }
```

```
    public int Stock { get; set; }
```

```
}
```

```
public class ProductContext : DbContext
```

```
{
```

```
    public ProductContext(DbContextOptions<ProductContext> options) : base(options) { }
```

```
    public DbSet<Product> Products { get; set; }
```

```
}
```

Register the context in `Program.cs`:

```
builder.Services.AddDbContext<ProductContext>(options =>  
    options.UseSqlServer(builder.Configuration.GetConnectionString("DefaultConnection")));
```

## Configure Elasticsearch

1. Create an `ElasticsearchService`:

using Nest;

public class ElasticsearchService

{ private readonly ElasticClient \_client;

public ElasticsearchService(IConfiguration configuration)

{

var settings = new ConnectionSettings(new Uri(configuration["Elasticsearch:Uri"]))

.DefaultIndex("products");

\_client = new ElasticClient(settings);

}

public async Task IndexProductAsync(Product product)

{

await \_client.IndexDocumentAsync(product);

}

```
public async Task<ISearchResponse<Product>> SearchProductsAsync(string query)
{
    return await _client.SearchAsync<Product>(s => s
        .Query(q => q
            .Match(m => m
                .Field(f => f.ProductName)
                .Query(query)
            )
        )
    );
}
```

Register the Elasticsearch service in `Program.cs`:

```
builder.Services.AddSingleton<ElasticsearchService>();
```

#### 4. Create API Endpoints

1. Add a `SearchController`:

```
[ApiController]
```

```
[Route("api/[controller]")]
```

```
public class SearchController : ControllerBase
```

```
{  
  
    private readonly ElasticsearchService _elasticService;  
  
    public SearchController(ElasticsearchService elasticService)  
    {  
        _elasticService = elasticService;  
    }  
}
```

```
[HttpPost("index")]
```

```
public async Task<ActionResult> IndexProduct([FromBody] Product product)
```

```
{
```

```
    await _elasticService.IndexProductAsync(product);
```

```
    return Ok("Product indexed successfully.");
```

```
}
```

```
[HttpGet("search")]
```

```
public async Task<ActionResult> Search([FromQuery] string query)
```

```
{
```

```
    var result = await _elasticService.SearchProductsAsync(query);
```

```
    return Ok(result.Documents);
```

```
}
```

```
}
```

## Testing

1. **Index a product into Elasticsearch:**

POST /api/search/index

Content-Type: application/json

```
{  
  "productID": 4,  
  "productName": "Monitor",  
  "price": 250,  
  "stock": 30  
}
```

**Search for products:**

GET /api/search/search?query=Monitor